

COMS 4771 Exam 2 Cheat Sheet

Elizabeth Kushelevsky, Spring 2025

Linear Algebra.....	2
Lp Norms.....	4
Decision Trees.....	5
Kernelization.....	8
Support Vector Machines.....	11

Linear Algebra

- Positive Definite: An n by n symmetric matrix A is PD if

$$\forall \mathbf{x} \in \mathbb{R}^n \setminus \{0\} : \mathbf{x}^T A \mathbf{x} > 0$$

- Translation: A is PD if all of its eigenvalues are strictly positive.
 - A has full rank and A is invertible.
 - There exists an invertible PD matrix B of the same rank as A such that $B = A^T A$.
- Positive Semidefinite: An n by n matrix symmetric A is PSD if

$$\forall \mathbf{x} \in \mathbb{R}^n : \mathbf{x}^T A \mathbf{x} \geq 0$$

- Translation: A is PSD if all of its eigenvalues are strictly nonnegative.
 - There exists a PSD matrix B of the same rank as A such that $B = A^T A$.
- Derivative of a linear function in vector form: for vectors a and w and scalar β ,

$$\frac{\partial}{\partial w} a^T w + \beta = a \quad (\text{from Jacobian})$$

- Derivative of a quadratic function in matrix form: for vector w and matrix A ,

$$\frac{\partial}{\partial w} w^T A w + \beta = (A^T + A)w \quad (\text{from Jacobian})$$

In the case that A is symmetric, like all PSD matrices, this simplifies to $2Aw$.

- Second derivative of a linear function in matrix form: for vectors a and w and scalar β ,

$$\frac{\partial^2}{\partial w^2} a^T w + \beta = 0 \quad (\text{from Hessian})$$

- Second derivative of a quadratic function in matrix form: for vector w and matrix A ,

$$\frac{\partial}{\partial w} w^T A w + \beta = A^T + A \quad (\text{from Hessian})$$

In the case that A is symmetric, like all PSD matrices, this simplifies to $2A$.

Let M be a symmetric positive semidefinite matrix.

1. Argue that for $\alpha > 0$, $M + \alpha I$ is positive definite. What are the eigenvalues w.r.t. the eigenvalues of M ?

New eigenvalues are $\lambda_i + \alpha$. If it is given that $\lambda_i \geq 0$ and $\alpha > 0$, then $\lambda_i + \alpha > 0$.

2. Find $\operatorname{argmin}_w \left[\frac{1}{2} w^T (M + \alpha I) w - b \cdot w \right]$ for an arbitrary vector b .

$$\frac{d}{dw} \left(\frac{1}{2} w^T (M + \alpha I) w - b \cdot w \right) = (M + \alpha I) w - b$$

$$(M + \alpha I) w - b = 0 \rightarrow (M + \alpha I) w = b$$

$$(M + \alpha I)^{-1} (M + \alpha I) w = (M + \alpha I)^{-1} b$$

$$w = (M + \alpha I)^{-1} b$$


3. Find $\max_w \left[\frac{1}{2} w^T (M + \alpha I) w - b \cdot w \right]$ for an arbitrary vector b .

Remember that this is a quadratic function shown as a PD matrix (convex set), which definitionally goes to ∞ . By choosing any w orthogonal to b , the $b \cdot w$ term drops out and the $\frac{1}{2} w^T (M + \alpha I) w$ term goes to infinity.


L_p Norms

- Normed linear space L_p : a pair of objects (V, d) where V is the chosen vector space and d is a norm (length measurement). $d: V \times V = ||x|| \rightarrow \mathbb{R}$
- L_p norms have a specific form corresponding to their p : $||x||_p = \sqrt[p]{\sum_{i=0}^n |x_i|^p}$, where n is the dimensionality of V .
 - Example: $||x||_3 = (x_1^3 + x_2^3 + \dots + x_n^3)^{1/3}$
- Properties of L_p norms:
 - $||x||_p \geq 0$
 - Only the zero vector in V has a norm of 0
 - $||\lambda x||_p = \lambda ||x||_p \forall \lambda \in \mathbb{R}$ (closed under scalar multiplication)
 - $||x_1 - x_2||_p \leq ||x_1||_p + ||x_2||_p$ (triangle inequality)
 - $p < 0$ does not define a norm because it fails this property. This means that the set of values $||x||_p \leq 1$ is NOT a convex set.
- L_0 is a special norm: $||x||_0 = \left(\sum_{i=0}^n |x_i|^0 \right)^\infty$ i.e. the number of nonzero components of x .
Can be used to assess the vector's sparsity.
- L_∞ is a special norm: $||x||_\infty = \max_i (|x_i|)$ i.e. the absolute value of the largest component of x .
- Common “unit spheres” of L_p norms (shown in \mathbb{R}^2):
 - L_1 : we are looking for vectors $x = (x_1, x_2)$ s.t. $||x||_1 = \sqrt[1]{x_1 + x_2} = 1$. Starting from the elementary vectors, these unit vectors form a diamond.
 - L_2 : same logic as before. Result is the unit circle.
 - $L_n, 2 < n < \infty$: approaches the unit sphere of L_∞ . A square with progressively pointier edges.
 - L_∞ : must have at least one component of length 1. Result is a square.

In \mathbb{R}^2 :



In \mathbb{R}^3 :



Decision Trees

- Idea: split data along decision boundaries, then classify by region. To split smartly, pick the label that maximally reduces label uncertainty (helps accurate classification the most)
- Ways to measure uncertainty:
 - Classification error: where p_y is the fraction of training data labeled y ,

$$u(C) = 1 - \max_y p_y$$
 - Entropy: corresponds to having 0 uncertainty when all info is known b/c of log.

$$u(C) = \sum_{y \in Y} p_y \log \frac{1}{p_y}$$

- Gini index:

$$u(C) = \sum_{y \in Y} p_y^2$$


Perceptron Algorithm

- Decision boundary is affine (form $w(x) = w \cdot x + w_0 = 0$)
- The linear classifier f originally has an offset: $\text{sign}(w \cdot x + x_0)$. To make the classifier pass through 0 and therefore closed under scalar multiplication, perform a “lifting” on it: $w' \cdot x' = (w, w_0) \cdot (x, 1)$
- Given training data $\{(x_1, y_1), \dots, (x_n, y_n)\}$, our goal is to find the weight vector w which minimizes training error:

$$\text{argmin}_w \frac{1}{n} \sum_{i=1}^n 1[\text{sign}(w \cdot x_i) \neq y_i] \quad (\text{Proportion of wrong classifications})$$

$$= \text{argmin}_w \left(\sum_{x_i \text{ s.t. } y_i = +1}^n 1[(x_i \cdot w) < 0] \right) + \left(\sum_{x_i \text{ s.t. } y_i = -1}^n 1[(x_i \cdot w) \geq 0] \right)$$

- Algorithm: under the assumption that the data IS linearly separable,
Initialize $w^{(0)} = 0$
For $t = \text{time index}$ {
 Propose a w and assess its training error.
 If $\exists (x, y)$ s.t. $\text{sign}(w^{(t-1)} \cdot x) \neq y$ (Misclassified point exists)
 $w^t = \begin{cases} w^{(t-1)} + x & \text{if } y = +1 \\ w^{(t-1)} - x & \text{if } y = -1 \end{cases}$
 } Terminate when no mistakes remain.
- Geometrically, perceptron algorithm relies on the cosine inequality:
 $w \cdot x = ||w|| ||x|| \cos \theta$
 - Because cosine is positive for angles above the proposed decision boundary, it creates positive and negative sides of the boundary.



Here, point x is being misclassified as negative, so w^t will adjust to classify it as positive.

- Perceptron mistake bound: for data which is linearly separable by a unit-length boundary w^* , there is an upper bound on the number of iterations the perceptron algorithm can

make before it perfectly separates the data, $\left(\frac{R}{\gamma}\right)^2$, where R is the radius of the data (distance to farthest point) and γ is the margin (distance to closest point).

- R in the numerator: when the adjustment is by a large vector, we expect w to have a larger “swing”, meaning more misclassifications
- γ in the denominator: when the margin between oppositely-signed points is closer, it takes more iterations to find a w with no error
- Proof: we will analyze the closeness between an arbitrary current w^t and w^* using the dot product. Suppose this is not the first weight vector tried, and therefore $w^{(t-1)}$ classified a point inaccurately.

$$w^t = (w^{(t-1)} + \gamma x)$$

Now, take the inner product with w^* .

$$w^t \cdot w^* = (w^{(t-1)} + \gamma x) \cdot w^*$$

Because γ is the closest point to the decision boundary, offsetting $w^{(t-1)}$ to w^t means shifting the boundary by at least γ . This way, we get a lower bound for the RHS of the equation.

$$\begin{aligned} (w^{(t-1)} + \gamma x) &\geq w^{(t-1)} + \gamma \\ (w^{(t-1)} + \gamma x) \cdot w^* &\geq w^{(t-1)} \cdot w^* + \gamma \\ w^t \cdot w^* &\geq w^{(t-1)} \cdot w^* + \gamma \end{aligned}$$

Now we look at the norm of w^t .

$$\begin{aligned} \|w^t\|^2 &= \|w^{(t-1)} + \gamma x\|^2 \\ &= \|w^{(t-1)}\|^2 + 2\gamma(w^{(t-1)} \cdot x) + \|\gamma x\|^2 \end{aligned}$$

The second term is always non-positive since the prediction was wrong, so we take it at its maximum, 0, and it disappears from the equation. The upper bound for the distance from the sign-adjusted misclassified point is the radius R .

$$\|w^t\|^2 \leq \|w^{(t-1)}\|^2 + R^2$$

This inequality and the γ inequality hold true for all iterations of the algorithm.

From the γ inequality, after T cycles of the algorithm, w must have shifted by at least $T \cdot \gamma$.

$$w^T \cdot w^* \geq w^0 \cdot w^* + T\gamma$$

Because w^0 is defined to be the 0 vector, its term goes away.

$$w^T \cdot w^* \geq T\gamma$$

By the Cauchy-Schwarz inequality ($a \cdot b \leq \|a\| \|b\|$),

$$w^T \cdot w^* \leq \|w^T\| \|w^*\| \quad (w^* \text{ is a unit vector!})$$

$$T\gamma \leq \|w^T\| \|w^*\|$$

From norm inequality, we have $\|w^T\|^2 \leq \|w^0\|^2 + R^2 T \rightarrow \|w^T\|^2 \leq R^2 T$.

$$\|w^T\| \leq R\sqrt{T} \quad (\text{which also bounds } T\gamma)$$

$$T \leq \left(\frac{R}{\gamma}\right)^2$$

■

Given n linearly independent feature vectors in n dimensions, show that for any assignment to the binary labels, you can always construct a linear classifier with weight vector w which separates the points. Assume that the classifier has the form $\text{sign}(w \cdot x)$. Note that a square matrix composed of linearly independent rows is invertible.

To make the sign function able to classify binary data, simply assign $y = \pm 1$ as the binary outcomes and fill in the features and dimensions into a matrix X . We claim that a weight vector w exists s.t. Every point is correctly classified, i.e. $\text{sign}(X * w) = y$, which is true iff $(X * w) = y$. Then, since X is invertible, $w = X^{-1}y$ and w exists for any features and assignments.

Kernelization

- Introduced as an extension of the perceptron algorithm into non-linearly separable data.
Theorem: all data is linearly separable in some space.
 - Shortcomings: data can be sent to an infinite-size vector where it can't be simply read, many points means a large computational cost, and the model cannot make a prediction on a posteriori data if it depends on individual a priori points
- Kernelization massages vectors so that their time-costly operations can be rewritten in terms of dot products, reducing computation time to $O(d)$. It *implicitly* considers vectors in higher dimensions through their similarity/dissimilarity without actually translating the vectors.
 - Method: rewrite your desired function using only dot products, then replace those dot products with a predetermined kernelized distance metric $k(x, z)$.

- Example: consider data that is quadratically separable in (x_1, x_2) . The linearization of points from $(x_1, x_2) \rightarrow (x_1^2, x_2^2)$ is a complicated operation:

$$g(x) = w_1 x_1^2 + w_2 x_2^2 + w_3 x_1 x_2 + w_4 x_1 + w_5 x_2 + w_0$$

$$= \sum_{i,j=1}^d \sum_{p+q \leq 2} w_{ij}^{p,q} x_i^p x_j^q$$

Instead, first define an explicit operation for $g(x)$ as the scaling of its coordinates:

$$\phi(x) = (x_1^2, \dots, x_d^2, \sqrt{2}x_1 x_2, \dots, \sqrt{2}x_{d-1} x_d, \sqrt{2}x_1, \dots, \sqrt{2}x_d, 1)$$

Then, the distance between two points $\phi(p), \phi(q)$ becomes computable as

$$\phi(p) \cdot \phi(q) = (p \cdot q)^d = K(p, q).$$

- Kernelizing the perceptron algorithm: here, the costly operation we hope to make computable in linear time is the calculation of w^t from $w^{t-1} \dots w^0$.

$$w = \sum_{k=1}^n \alpha_k y_k x_k \quad (\alpha = \text{number of mistakes on } x_k)$$

Instead, look at how data is classified based on this weighting:

$$f(x) = \text{sign}(w \cdot x)$$

$$= \text{sign}\left(\left(\sum_{k=1}^n \alpha_k y_k x_k\right) \cdot x\right) = \text{sign}\left(\sum_{k=1}^n \alpha_k y_k (x_k \cdot x)\right)$$

In any non-linearly separable space which we hope to linearize, instead of carrying out the transformation $\phi(x)$, we can use the kernelized dot product.

$$f(x) = \text{sign}\left(\sum_{k=1}^n \alpha_k y_k (k(x_k, x))\right) \quad (\text{For the user's choice of } k)$$

- Some more common kernels:

Name	Kernel Function (implicit dot product)	Feature Space (explicit dot product)
Linear	$K(\mathbf{x}, \mathbf{z}) = \mathbf{x}^T \mathbf{z}$	Same as original input space
Polynomial (v1)	$K(\mathbf{x}, \mathbf{z}) = (\mathbf{x}^T \mathbf{z})^d$	All polynomials of degree d
Polynomial (v2)	$K(\mathbf{x}, \mathbf{z}) = (\mathbf{x}^T \mathbf{z} + 1)^d$	All polynomials up to degree d
Gaussian	$K(\mathbf{x}, \mathbf{z}) = \exp\left(-\frac{\ \mathbf{x} - \mathbf{z}\ _2^2}{2\sigma^2}\right)$	Infinite dimensional space
Hyperbolic Tangent (Sigmoid) Kernel	$K(\mathbf{x}, \mathbf{z}) = \tanh(\alpha \mathbf{x}^T \mathbf{z} + c)$	(With SVM, this is equivalent to a 2-layer neural network)

Support Vector Machines

- General idea: find the *optimal* decision boundary as the midpoint of the two maximally wide boundaries. Minimize the inverse of this distance, $\frac{1}{2}||w^2||$, subject to the constraint for minimum clearance, $y_i(w_i \cdot x_i - b) \geq 1$.

- If we wish to allow slack, the problem becomes $\min \frac{1}{2}||w^2|| + C \sum_{i=1}^n \xi_i$ subject to $y_i(w_i \cdot x_i - b) \geq 1 - \xi_i$, where ξ_i is the slack allowed on point i and C is a user-controlled hyperparameter accounting for the tradeoff between slack and mistakes. Large $C \rightarrow$ assign large penalty to mistake

- SVM is a convex optimization problem, where a convex objective function f is minimized over a convex feasible region subject to constraints $g_i(x) \leq 0$.

- It can be transformed into the Lagrangian function $L(x, \lambda)$, where λ is a scaling vector on the constraints.

$$L(x, \lambda) = f(x) + \sum_{i=1}^n \lambda_i g_i(x)$$

- Given the primal $p^* = \min_x \max_{\lambda} L(x, \lambda)$ and the dual $d^* = \max_{\lambda} \min_x L(x, \lambda)$, If the constraints can be strict (Slater's condition: $g_i(x) < 0 \forall g$ OR $g_i(x) \leq 0$ where g is affine $\rightarrow g$ creates a convex feasible region), $p^* = d^*$ (strong Lagrangian duality)

- For SVM, we have the following:

Minimize w, b for $f(w) = \frac{1}{2}||w^2||$

With constraints $y_i(w_i \cdot x_i - b) \geq 1$

The Lagrangian is:

$$L(w, b, \alpha) = \frac{1}{2}||w^2|| + \sum_{i=1}^n \alpha_i (y_i(w_i \cdot x_i - b))$$

and $\min_{w, b}$ can be found by taking the gradient of L .

$$\frac{\partial}{\partial w} L(w, b, \alpha) = w - \sum_{i=1}^n \alpha_i y_i x_i = 0$$

$$w = \sum_{i=1}^n \alpha_i y_i x_i$$

$$||w^2|| = w^T w = \left(\sum_{i=1}^n \alpha_i y_i x_i \right)^T \left(\sum_{j=1}^n \alpha_j y_j x_j \right)$$

$$= \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \alpha_i \alpha_j y_i y_j (x_i \cdot x_j)$$

Now we have the SVM standard form.

$\text{Maximize } \alpha_i \text{ on } \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \alpha_i \alpha_j y_i y_j (x_i \cdot x_j)$ $\text{such that } \sum_{i=1}^n \alpha_i y_i, \alpha_i \geq 0$

- Because this maximization only performs dot products on x , it is kernelized.